

EECE6036 - Homework 3

Wayne Stegner

October 31, 2020

1 Problem 1

1.1 Problem Summary

The goal of this problem is to create a multi-layer feed-forward neural network for classification of the MNIST dataset. The network was trained using gradient descent with back-propagation with momentum, and is trainable for arbitrary amounts and sizes of hidden layers. During training, the loss is calculated using J_2 loss with threshold values. For performance measurement, the loss is calculated as $1 - \text{hit_rate}$ using winner-take-all on the output layer to define the predicted class.

1.2 Results

1.2.1 System Description

Table 1 shows the hyper-parameters used in training the classifier. These hyper-parameters were found empirically, considering both minimizing the final loss and training time. More work can to find more optimal hyper-parameters, but this set of hyper-parameters produces pretty good results.

Table 1: Classifier Training Hyper-Parameters

Parameter	Value	Description
<i>hidden_layer_size</i>	192	Neurons in hidden layer
η	0.05	Learning rate
α	0.8	Momentum
<i>max_epochs</i>	500	Maximum training epochs
L	0.25	Lower activation threshold
H	0.75	Upper activation threshold
<i>patience</i>	3	Patience before early stopping
<i>es_delta</i>	0.01	Delta value for early stopping

Weight initialization is done randomly on a uniform distribution between $(-a, +a)$ where $a = \sqrt{\frac{6}{N_{source} + N_{target}}}$, and N_{source} and N_{target} are the numbers of neurons on the source and target layers respectively.

The network utilizes early stopping by monitoring a validation set, which consists of 1000 training points that are set aside before training. If the validation loss does not improve for *patience* steps (1 step = 10 epochs), training halts. The validation is considered improved if the validation is *es_delta* lower than the previous improved validation loss. The weights used by the final network are the weights from the epoch with the lowest validation loss. To improve the frequency of checking for early stopping, the network only uses 1000 training points per epoch.

1.2.2 Network Results

Throughout the duration of training, the loss of the training and validation sets was tracked every 10 epochs. Figure 1 shows a plot of the loss values while training the classifier. The vertical line designates the point where the validation error is minimized, which occurs at epoch 50. The weights at that epoch are the weights used in the final network, where the test loss is 0.070.




Figure 1: Training and validation loss of the classifier.

Figure 2 shows the confusion matrices for the train and test sets using the weights from the best epoch.




Figure 2: Confusion matrices of the train and test sets.

1.3 Discussion and Analysis of Results

Overall, the results look pretty good. The best weights happen pretty early at epoch 50 and get a fairly low loss of 0.070. While testing hyper-parameters, it was observed that when *es_delta* is set to 0 and *patience* is increased to around 5, the test loss would sometimes get as low as 0.05, but that generally occurred after several hundred epochs and is also dependent on weight initialization numbers and dataset shuffling during training. With the early stopping hyper-parameters used by this model, it sacrifices some performance, but it trains significantly faster, making these hyper-parameters ideal for implementing and troubleshooting the network from scratch.

The diagonal of the confusion matrices in Figure 2 is very dark, with hardly any coloring in the incorrect boxes. Many of the incorrect classifications are in two classes which look similar. For example, Figure 2a shows four 9s classified as 4s and five 4s classified as 9s. This mix-up makes a lot of sense, because the shapes of 4s and 9s tend to be similar, especially with some of the sloppy handwriting in MNIST.

1.4 Conclusion

This multi-layer feed-forward neural network classifier is able to effectively classify the MNIST dataset given in this problem. While the results are certainly not state-of-the-art, the trained network classifies over 90% of the test set digits correctly. Further testing in optimizing the hyper-parameters of the network can help to improve the accuracy even further.

2 Problem 2

2.1 Problem Summary

The goal of this problem is to create a multi-layer feed-forward neural network similar to Problem 1, except in this case it will be trained to be an autoencoder. This network has the same features as the network in Problem 1, except the performance measurement is calculated using J_2 loss instead of $1 - hit_rate$. Additionally, winner-take-all is not used in the output layer, because it does not make sense in the context of an autoencoder.

2.2 Results

2.2.1 System Description

Table 2 shows the hyper-parameters used in training the autoencoder. As with the classifier, these hyper-parameters were found empirically with the same consideration for minimizing final loss and training time.

Table 2: Autoencoder Training Hyper-Parameters

Parameter	Value	Description
<i>hidden_layer_size</i>	192	Neurons in hidden layer
η	0.005	Learning rate
α	0.8	Momentum
<i>max_epochs</i>	500	Maximum training epochs
L	0	Lower activation threshold
H	1	Upper activation threshold
<i>patience</i>	3	Patience before early stopping
<i>es_delta</i>	0.1	Delta value for early stopping

The weight initialization process is the same as in Problem 1, as are the processes for allocating the validation set and applying early stopping.

2.2.2 Network Results

Throughout the duration of training, the loss of the training and validation sets was tracked every 10 epochs. Figure 3 shows a plot of the loss values while training the autoencoder. The vertical line designates the point where the validation error is minimized, which occurs at epoch 160. The weights at that epoch are the weights used in the final network, where the test loss is 2.087.

After training, the loss in each class was measured. Figure 4 shows the loss of each class for the train and test sets using the weights from the best epoch.




Figure 3: Training and validation loss of the autoencoder.




Figure 4: Loss of each class of the train and test sets.

2.2.3 Features

After training, the features learned by 20 arbitrary neurons in the hidden layer of both the classifier and autoencoder were observed by normalizing the weights from 0 to 1 and then displaying them like an image. The images are mapped so 1 is white and 0 is black. Figure 5 shows the features of the classifier, and Figure 6 shows the features of the autoencoder.

2.2.4 Sample Outputs

After training, the outputs of eight random data points from the test set were fed reconstructed the autoencoder, meaning the input data point was presented to the network and the output "prediction" was obtained. Figure 7 shows the original and reconstructed images.




Figure 5: Classifier features.




Figure 6: Autoencoder features.

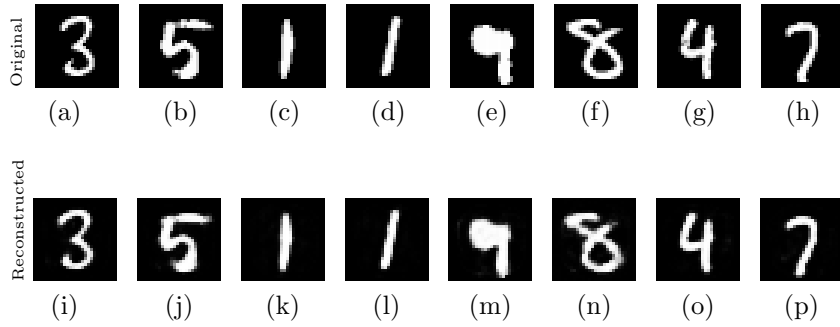


Figure 7: Original (top) and reconstructed (bottom) data points.

2.3 Discussion and Analysis of Results

Overall, the results look pretty good. The autoencoder takes longer to train, with its best weights occurring at epoch 160 with a loss of 2.087. In Figure 4, the loss for the 1s is particularly salient because it is drastically lower than the other classes loss values. This makes sense because a 1 is simply a straight line, so it should be quite simple to recognize and then reconstruct. While none of the loss values are as noticeably high as 1's loss is low, classes 2, 8, and 5 all have loss values above 2.5. I do not have any intuition as to why those numbers have the highest loss, but those digits are kind of similar in the sense that a 2 is roughly a backward 5, and stacking a 2 on top of a 5 roughly makes an 8. At first, I was surprised that 6 has a higher loss than

9, because 6 is just an upside-down 9, but it makes sense that 5 and 6 have similar features and ended up having similar loss values.

The features in Figure 5 and Figure 6 are somewhat difficult to interpret. Between the two feature sets, there is a trend where the center of the image appears to have some distinct patterns, while the edge of the image is either somewhat uniform gray or random TV static. In Figure 5, the shapes generally seem to be lines or curves, while in Figure 6 the features look like a bunch of dark dots with a few brighter spots.

2.4 Conclusion

This multi-layer feed-forward neural network autoencoder is able to effectively reconstruct digits from the MNIST dataset. Further testing in optimizing the hyper-parameters of the network can help to improve the reconstruction accuracy even further.

A Code

A.1 dataset.py

```
#####  
# Imports  
#####  
from settings import *  
import numpy as np  
import pathlib  
import matplotlib.pyplot as plt  
  
def shufflePair(data, labels):  
    '''Shuffle a pair of data and labels in place  
    Parameters:  
    -----  
        data, labels : np.ndarray  
        Data and labels to be shuffled  
    Returns:  
    -----  
        None  
    '''  
    assert len(data) == len(labels), \  
        'Size mismatch between data and labels'  
    indeces = np.random.permutation(len(data))  
    data[...] = data[indeces]  
    labels[...] = labels[indeces]  
  
if __name__ == '__main__':  
    # Define important values from settings  
    SEED = 69420  
    CODE_DIR = pathlib.Path(__file__).parent.absolute()  
    DATA_FILE = CODE_DIR.joinpath('data.txt')  
    LABEL_FILE = CODE_DIR.joinpath('labels.txt')  
    DATASET_DIR = CODE_DIR.joinpath('dataset')  
    DATASET_DIR.mkdir(mode=0o755, exist_ok=True)  
    TRAIN_DATA_FILE = DATASET_DIR.joinpath('train_data.npy')  
    TRAIN_LABELS_FILE = DATASET_DIR.joinpath('train_labels.npy')  
    TEST_DATA_FILE = DATASET_DIR.joinpath('test_data.npy')  
    TEST_LABELS_FILE = DATASET_DIR.joinpath('test_labels.npy')  
    # Seed for consistency  
    np.random.seed(SEED)  
    #####  
    # Import dataset  
    #####  
    # Check that the files exist  
    assert DATA_FILE.exists(), f'File not found: {str(data_file)}'  
    assert LABEL_FILE.exists(), f'File not found: {str(label_file)}'  
    # Read data and labels from txt file  
    data = list()  
    with open(DATA_FILE) as data_f:  
        data = data_f.readlines()  
    data = [d.split() for d in data]  
    labels = list()  
    with open(LABEL_FILE) as label_f:  
        labels = label_f.readlines()  
    labels = [int(l) for l in labels]  
    # Make images and sort into bins based on class  
    data_points = dict()  
    classes = list()  
    for d, l in zip(data, labels):  
        if l not in classes:  
            classes.append(l)  
            data_points[l] = list()  
            data_points[l].append(d)  
    classes.sort()  
    # Turn data lists into numpy arrays  
    for l in classes:  
        data_points[l] = np.array(data_points[l], dtype=np.float64)  
    #####  
    # Partition dataset  
    #####  
    train_data = list()  
    train_labels = list()  
    test_data = list()  
    test_labels = list()  
    # Shuffle the data points first  
    for l in classes:  
        np.random.shuffle(data_points[l])
```

```

# Iterate through the data by class and partition
points_per_class = len(data_points[classes[0]])
train_per_class = int(TRAIN_PORTION * points_per_class)
for l in classes:
    for i, d in enumerate(data_points[l]):
        if i < train_per_class:
            train_data.append(d)
            train_labels.append(1)
        else:
            test_data.append(d)
            test_labels.append(1)
# Turn lists into numpy arrays
train_data = np.array(train_data)
train_labels = np.array(train_labels)
test_data = np.array(test_data)
test_labels = np.array(test_labels)
# Turn labels into one-hot arrays
train_labels = np.eye(len(classes))[train_labels]
test_labels = np.eye(len(classes))[test_labels]
# Shuffle arrays
shufflePair(train_data, train_labels)
shufflePair(test_data, test_labels)
#####
# Save arrays
#####
np.save(str(TRAIN_DATA_FILE), train_data)
np.save(str(TRAIN_LABELS_FILE), train_labels)
np.save(str(TEST_DATA_FILE), test_data)
np.save(str(TEST_LABELS_FILE), test_labels)

```

A.2 layer.py

```

#####
# Imports
#####
import numpy as np

class Layer:
    def __init__(
        self,
        weight_file = None,
        num_neurons = None,
        inputs = None,
    ):
        '''Initialize Layer object either randomly or by a weight file
        Parameters:
        -----
            weight_file : str, optional
                File to load pre-existing weights from
            num_neurons, inputs : int, optional
                Number of hidden neurons and inputs to the layer
        Returns:
        -----
            Layer
            The Layer object which was constructed
        '''
        if weight_file:
            self.loadWeights(weight_file)
            self.num_neurons = self.w.shape[0]
            self.inputs = self.w.shape[1] - 1
        else:
            self.num_neurons = num_neurons
            self.inputs = inputs
            self.initW()
        # States to save for back-prop
        self.x = None # Input for current pass
        self.s = None # Net inputs pre-activation function
        self.y = None # Final output of the layer
        self.delta = None # Delta of this layer
        self.w_change = np.zeros(self.w.shape) # Weight changes

    def initW(self):
        '''Initialize weights for the perceptron using Xavier initialization
        Parameters:
        -----
            None
        Returns:
        -----
            None
        '''
        w_shape = (self.num_neurons, self.inputs + 1)

```

```

a = np.sqrt(6 / (self.inputs + self.num_neurons))
self.w = np.random.uniform(-a, a, size=w_shape)
self.w_change = np.zeros(self.w.shape)

def forwardPass(self, x):
    '''Pass the input forward through the layer
    Parameters:
    -----
        x : np.ndarray of np.float64
            Input to the layer
    Returns:
    -----
        np.array of np.float64
            Output of the layer
    '''
    # Add bias input
    x = np.concatenate([[1], x])
    # Save inputs for back-prop
    self.x = x
    # Dot product weights * inputs
    self.s = np.matmul(self.w, x)
    # Pass through sigmoid activation
    self.y = 1.0 / (1 + np.exp(-1 * self.s))
    return self.y

def setDelta(self):
    '''Calculate delta value, different for Output and Hidden layer
    Must be implemented per-class
    '''
    raise NotImplementedError('Cannot call from Layer class')

def getWChange(self, eta=1, alpha=0.1):
    '''Calculate weight updates for the most recent forward pass
    Requires delta to be calculated (varies between hidden/output layers)
    Parameters:
    -----
        eta : float
            Learning rate for weight adjustments
        alpha : float
            Momentum scalar
    Returns:
    -----
        None
    '''
    # Set delta value
    self.setDelta()
    # Pre-scale weight change for momentum
    self.w_change *= alpha
    # Calculate new weight change
    d = self.delta.reshape(len(self.delta), 1)
    x = self.x.reshape(1, len(self.x))
    new_change = eta * np.matmul(d, x)
    self.w_change += new_change

def changeW(self):
    '''Apply w_change to the weights
    Parameters:
    -----
        None
    Returns:
    -----
        None
    '''
    self.w += self.w_change

def saveWeights(self, weight_file):
    '''Save weights to a file
    Parameters:
    -----
        weight_file : str
            File name to save weights in
    Returns:
    -----
        None
    '''
    np.save(str(weight_file), self.w)

def loadWeights(self, weight_file):
    '''Save weights to a file
    Parameters:
    -----
        weight_file : str
            File name to save weights in

```

```

        Returns:
        -----
        None
    '''
    self.w = np.load(str(weight_file))

if __name__ == '__main__':
    print('Warning: Tests for this file are deprecated')

```

A.3 output_layer.py

```

#####
# Imports
#####
from layer import Layer
import numpy as np

class OutputLayer(Layer):
    def setLabel(self, label):
        '''Set the labels for this batch
        Needed for calculating delta for this layer
        Parameters:
        -----
            label : np.ndarray
                   One-hot encoded label
        Returns:
        -----
            None
        '''
        self.label = label

    def setDelta(self):
        '''Calculate delta for the output layer
        Parameters:
        -----
            None
        Returns:
        -----
            None
        '''
        e = self.label - self.y
        y_der = self.y * (1 - self.y)
        self.delta = e * y_der

    def thresholdOutputs(self, L, H):
        '''Threshold outputs based on the labels
        Parameters:
        -----
            L, H: float
                  Low and high thresholds for outputs
        Returns:
        -----
            None
        '''
        self.y[(self.y <= L) * (self.label == 0)] = 0
        self.y[(self.y >= H) * (self.label == 1)] = 1

if __name__ == '__main__':
    print('Warning: Tests for this file are deprecated')

```

A.4 hidden_layer.py

```

#####
# Imports
#####
from layer import Layer
import numpy as np

class HiddenLayer(Layer):
    def setDownstreamSum(self, w, delta):
        '''Sum the product of w and delta for the next layer
        Needed for calculating delta for this layer
        Parameters:
        -----
            w : np.ndarray

```

```

        Matrix of weight values for the next layer
        delta : np.ndarray
        Matrix of delta values for the next layer
Returns:
-----
    None
'''
self.downstream_sum = np.matmul(w[:, :-1].transpose(), delta)

def setDelta(self):
    '''Calculate delta for the hidden layer
    Parameters:
    -----
        None
    Returns:
    -----
        None
    '''
    # Derivative of sigmoid using last forward pass
    output_der = self.y * (1 - self.y)
    self.delta = output_der * self.downstream_sum

if __name__ == '__main__':
    print('Warning: Tests for this file are deprecated')

```

A.5 mlp.py

```

#####
# Imports
#####
# Custom imports
from settings import *
from dataset import shufflePair
from hidden_layer import HiddenLayer
from output_layer import OutputLayer
# External imports
import numpy as np
import enlighten # Progress bar for training

class MLP:
    def __init__(
        self,
        input_size=INPUTS,
    ):
        '''Initialize Multi-Layer Perceptron object
        Parameters:
        -----
            input_size : int
                Number of inputs to the network
        '''
        # Initialize parameters
        self.layers = list()
        self.input_size = input_size

    def addLayer(self, file_name=None, neurons=None, output=False):
        '''Add a layer to the network
        '''
        if len(self.layers) == 0:
            input_size = self.input_size
        else:
            input_size = self.layers[-1].num_neurons
        if file_name:
            if output:
                self.layers.append(OutputLayer(weight_file=file_name))
            else:
                self.layers.append(HiddenLayer(weight_file=file_name))
        else:
            if output:
                self.layers.append(OutputLayer(num_neurons=neurons,
                    inputs=input_size))
            else:
                self.layers.append(HiddenLayer(num_neurons=neurons,
                    inputs=input_size))

    def predict(self, data, one_hot):
        '''Predict the output given an input
        Parameters:
        -----
            data : np.ndarray
                Data point to predict on

```

```

        one_hot : bool
            Whether the output should be one-hot or raw
Returns:
-----
    np.ndarray
        Array of prediction
    """
    # Forward pass through all the layers
    layer_output = data
    for l in self.layers:
        layer_output = l.forwardPass(layer_output)
    if one_hot:
        max_idx = np.argmax(layer_output)
        pred = np.eye(10)[max_idx]
    else:
        pred = layer_output
    return pred

def trainPoint(self, data, label, eta, alpha, L, H):
    """Update the weights for a single point
    Parameters:
    -----
        data : np.ndarray
            The data point
        labels : np.ndarray
            One-hot encoded label
        eta : float
            Learning rate
        alpha : float
            Momentum scalar
        L, H : float, optional
            Low and high thresholds for training
    """
    # Weight change for last layer
    self.predict(data, one_hot=False)
    self.layers[-1].setLabel(label)
    self.layers[-1].thresholdOutputs(L, H)
    self.layers[-1].getWChange(eta, alpha)
    # Back-prop error
    for i in range(len(self.layers)-2, -1, -1):
        down_w = self.layers[i+1].w
        down_delta = self.layers[i+1].delta
        self.layers[i].setDownstreamSum(down_w, down_delta)
        self.layers[i].getWChange(eta, alpha)
    # Apply weight changes
    for l in self.layers:
        l.changeW()

def logError(self,
             train_data,
             train_labels,
             valid_data,
             valid_labels,
             epoch_num,
             ):
    """Log error for train/test data for a given epoch
    Parameters:
    -----
        train_data : np.ndarray
            Array of data points for the train set
        train_labels : np.ndarray
            Array of labels for the train set
        valid_data : np.ndarray
            Array of data points for the validation set
        valid_labels : np.ndarray
            Array of labels for the validation set
        epoch_num : int
            The current epoch number
    """
    # Evaluate and log train error
    train_err = self.eval(train_data, train_labels)
    self.train_err.append(train_err)
    # Evaluate and log test error
    valid_err = self.eval(valid_data, valid_labels)
    self.valid_err.append(valid_err)
    # Record epoch number
    self.epoch_num.append(epoch_num)
    # Print out metrics
    print(f'Epoch {epoch_num}')
    print(f'\tTrain Loss:\t\t{train_err:0.03f}')
    print(f'\tValidation Loss:\t{valid_err:0.03f}')

def train(

```

```

        self,
        data,
        labels,
        points_per_epoch,
        valid_points,
        max_epochs,
        eta,
        alpha,
        L,
        H,
        patience,
        es_delta,
        save_dir,
    ):
        '''Train the network up to the desired number of epochs
Parameters:
-----
    data : np.ndarray
        Array of training data points
    labels : np.ndarray
        Array of training labels
    points_per_epoch : int
        Number of training points to use in each epoch
    valid_points : int
        Number of training points to set aside for validation
        These points are set aside before training
    max_epochs : int
        Maximum number of epochs to train
    eta : float
        Learning rate
    alpha : float
        Momentum scalar
    L, H : float
        Low and high thresholds for training
    patience : int
        Amount of epochs with no improvement before early stopping
    es_delta : float
        Required improvement for early stopping
    save_dir : pathlib.Path or str
        Directory to save best model parameters in
'''
    # Initialize progress bar
    pbar_manager = enlighten.get_manager()
    pbar = pbar_manager.counter(total=max_epochs, desc='Training',
                                unit='epochs', leave=False)
    # Lists to track metrics
    self.train_err = list()
    self.valid_err = list()
    self.epoch_num = list()
    # Set aside validation partition (after shuffling)
    assert valid_points + points_per_epoch <= len(data), \
        'Not enough data for validation points and points per epoch'
    shufflePair(data, labels)
    valid_data = data[:valid_points]
    valid_labels = labels[:valid_points]
    train_data = data[valid_points:]
    train_labels = labels[valid_points:]
    # Log initial accuracy (using whole train partition)
    self.logError(
        train_data,
        train_labels,
        valid_data,
        valid_labels,
        0,
    )
    # Iterate through epochs or until early stopping
    impatience = 0
    self.best_weights_epoch = 0
    best_loss = np.inf
    for e in range(1, max_epochs+1):
        pbar.update()
        shufflePair(train_data, train_labels)
        epoch_train_data = train_data[:points_per_epoch]
        epoch_train_labels = train_labels[:points_per_epoch]
        for d, l in zip(epoch_train_data, epoch_train_labels):
            self.trainPoint(d, l, eta, alpha, L, H)
        # Log data every 10 epochs
        if (e % 10) == 0:
            self.logError(
                epoch_train_data,
                epoch_train_labels,
                valid_data,
                valid_labels,

```

```

        e,
    )
    # Check for early stopping
    if ((best_loss - self.valid_err[-1]) >= es_delta):
        impatience = 0
        self.best_weights_epoch = e
        best_loss = self.valid_err[-1]
        for i, l in enumerate(self.layers):
            layer_name = save_dir.joinpath(f'layer_{i:02d}')
            l.saveWeights(layer_name)
    else:
        impatience += 1
        # We have become too impatient
        if impatience >= patience:
            print(f'* * * Early stopping hit * * *')
            break
pbar.close()

if __name__ == '__main__':
    print('Warning: This file does not do anything.')
    print('Run either classifier.py or autoencoder.py.')

```

A.6 classifier.py

```

#####
# Imports
#####
# Custom imports
from mlp import MLP
# External imports
import numpy as np
import matplotlib.pyplot as plt

class Classifier(MLP):
    def eval(self, data, labels):
        '''Evaluate error on a data set
        Parameters:
        -----
            data : np.ndarray
                Array of data points
            labels : np.ndarray
                Array of labels for the data
        Returns:
        -----
            float
                Error (1 - Hit Rate)
        '''
        correct = 0
        for d, l in zip(data, labels):
            # Make prediction
            pred = self.predict(d, one_hot=True)
            # Check equality with label
            if (np.sum(np.abs(pred - l)) == 0):
                correct += 1
        hit_rate = correct / len(data)
        return (1 - hit_rate)

if __name__ == '__main__':
    # Additional imports
    from settings import *
    import pathlib
    import csv
    # Seed for consistency
    np.random.seed(SEED)
    # Delete old model
    for f in CLASS_MODEL_DIR.iterdir():
        f.unlink()
    # Test network
    classifier = Classifier(input_size=INPUTS)
    for h in HIDDEN_LAYER_SIZES:
        classifier.addLayer(neurons=h, output=False)
    classifier.addLayer(neurons=CLASSES, output=True)
    # Train the network
    print('* * * Begin training classifier * * *')
    classifier.train(
        train_data,
        train_labels,
        points_per_epoch=CLASS_POINTS_PER_EPOCH,
        valid_points=CLASS_VALID_POINTS,

```



```

max_epochs=CLASS_MAX_EPOCHS,
eta=CLASS_ETA,
alpha=CLASS_ALPHA,
L=CLASS_L,
H=CLASS_H,
patience=CLASS_PATIENCE,
es_delta=CLASS_ES_DELTA,
save_dir=CLASS_MODEL_DIR,
)

# Plot loss over epochs
plt.figure()
plt.plot(classifier.epoch_num, classifier.train_err)
plt.plot(classifier.epoch_num, classifier.valid_err)
plt.axvline(x=classifier.best_weights_epoch, c='g')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.title(f'Classifier Loss vs Epoch Number')
plt.xlabel('Epoch number')
plt.xlim([0, classifier.epoch_num[-1]])
plt.ylabel('Loss')
plt.ylim([0, max(max(classifier.train_err), max(classifier.valid_err))])
plt.savefig(str(CLASS_LOSS_PLOT), bbox_inches='tight', pad_inches=0)
plt.close()

# Save parameters to CSV
csv_rows = list()
csv_rows.append(['Parameter', 'Value', 'Description'])
csv_rows.append(['$hidden\\_layer\\_size$', str(HIDDEN_LAYER_SIZES[0]),
'Neurons in hidden layer'])
csv_rows.append(['$\\eta$', str(CLASS_ETA), 'Learning rate'])
csv_rows.append(['$\\alpha$', str(CLASS_ALPHA), 'Momentum'])
csv_rows.append(['$max\\_epochs$', str(CLASS_MAX_EPOCHS),
'Maximum training epochs'])
csv_rows.append(['$L$', str(CLASS_L), 'Lower activation threshold'])
csv_rows.append(['$H$', str(CLASS_H), 'Upper activation threshold'])
csv_rows.append(['$patience$', str(CLASS_PATIENCE),
'Patience before early stopping'])
csv_rows.append(['$es\\_delta$', str(CLASS_ES_DELTA),
'Delta value for early stopping'])
with open(str(CLASS_PARAM_CSV), 'w') as csv_file:
    csv_writer = csv.writer(csv_file)
    csv_writer.writerows(csv_rows)

# Write best epoch
with open(str(CLASS_BEST_EPOCH), 'w') as best_epoch_file:
    best_epoch_file.write(str(classifier.best_weights_epoch))

```

A.7 autoencoder.py

```

#####
# Imports
#####
from mlp import MLP
import numpy as np
import matplotlib.pyplot as plt

class Autoencoder(MLP):
    def eval(self, data, labels):
        '''Evaluate error on a data set
        Parameters:
        -----
            data : np.ndarray
                Array of data points
            labels : np.ndarray
                Array of labels for the data
        Returns:
        -----
            float
                Error (1 - Hit Rate)
        '''
        # Calculate loss
        total_loss = 0
        for d, l in zip(data, labels):
            pred = self.predict(d, one_hot=False)
            loss = np.sum(np.square(l - pred))
            total_loss += loss / 2
        total_loss /= len(data)
        return total_loss

if __name__ == '__main__':
    # Additional imports
    from settings import *
    import pathlib

```

```

import csv
# Seed for consistency
np.random.seed(SEED)
# Delete old model
for f in AUTO_MODEL_DIR.iterdir():
    f.unlink()
# Training constants
# Test network
autoencoder = Autoencoder(input_size=INPUTS)
for h in HIDDEN_LAYER_SIZES:
    autoencoder.addLayer(neurons=h, output=False)
autoencoder.addLayer(neurons=INPUTS, output=True)
# Train the network
print('* * * Begin training autoencoder * * *')
autoencoder.train(
    train_data,
    train_data,
    points_per_epoch=AUTO_POINTS_PER_EPOCH,
    valid_points=AUTO_VALID_POINTS,
    max_epochs=AUTO_MAX_EPOCHS,
    eta=AUTO_ETA,
    alpha=AUTO_ALPHA,
    L=AUTO_L,
    H=AUTO_H,
    patience=AUTO_PATIENCE,
    es_delta=AUTO_ES_DELTA,
    save_dir=AUTO_MODEL_DIR,
)
plt.figure()
plt.plot(autoencoder.epoch_num, autoencoder.train_err)
plt.plot(autoencoder.epoch_num, autoencoder.valid_err)
plt.axvline(x=autoencoder.best_weights_epoch, c='g')
plt.legend(['Train', 'Validation'], loc='upper right')
plt.title(f'Autoencoder Loss vs Epoch Number')
plt.xlabel('Epoch number')
plt.xlim([0, autoencoder.epoch_num[-1]])
plt.ylabel('Loss')
plt.ylim([0, max(max(autoencoder.train_err), max(autoencoder.valid_err))])
plt.savefig(str(AUTO_LOSS_PLOT), bbox_inches='tight', pad_inches=0)
plt.close()
# Save parameters to CSV
print('Saving training parameters')
csv_rows = list()
csv_rows.append(['Parameter', 'Value', 'Description'])
csv_rows.append(['$hidden\\_layer\\_size$', str(HIDDEN_LAYER_SIZES[0]),
    'Neurons in hidden layer'])
csv_rows.append(['$\\eta$', str(AUTO_ETA), 'Learning rate'])
csv_rows.append(['$\\alpha$', str(AUTO_ALPHA), 'Momentum'])
csv_rows.append(['$max\\_epochs$', str(AUTO_MAX_EPOCHS),
    'Maximum training epochs'])
csv_rows.append(['$L$', str(AUTO_L), 'Lower activation threshold'])
csv_rows.append(['$H$', str(AUTO_H), 'Upper activation threshold'])
csv_rows.append(['$patience$', str(AUTO_PATIENCE),
    'Patience before early stopping'])
csv_rows.append(['$es\\_delta$', str(AUTO_ES_DELTA),
    'Delta value for early stopping'])
with open(str(AUTO_PARAM_CSV), 'w') as csv_file:
    csv_writer = csv.writer(csv_file)
    csv_writer.writerows(csv_rows)
# Write best epoch
with open(str(AUTO_BEST_EPOCH), 'w') as best_epoch_file:
    best_epoch_file.write(str(autoencoder.best_weights_epoch))

```

A.8 test_classifier.py

```

#####
# Imports
#####
# Custom imports
from settings import *
from classifier import Classifier
# External imports
import numpy as np
import pathlib
import matplotlib.pyplot as plt

def makeConfMat(classifier, data, labels, plot_name, title):
    '''Generate and save a confusion matrix
    Parameters:
    -----
        data : np.ndarray

```

```

        Array of data values
    labels : np.ndarray
        Array of labels as one-hot-vectors
    plot_name : pathlib.Path or str
        File name to save the matrix as
    title : str
        Title of the confusion matrix
Returns:
-----
    None
'''
assert len(data) == len(labels), \
    'Size mismatch between data and labels'
conf_mat = np.zeros((CLASSES, CLASSES))
for d, l in zip(data, labels):
    pred = classifier.predict(d, one_hot=True)
    conf_mat += l.reshape((CLASSES,1)) * pred.reshape((1,CLASSES))
# Plot confusion matrix and save
plt.figure()
plt.suptitle(title)
plt.imshow(conf_mat, cmap='Greens')
for i in range(len(conf_mat)):
    for j in range(len(conf_mat[i])):
        color = 'k' if (conf_mat[i][j] <= 50) else 'w'
        plt.text(j, i, f'{int(conf_mat[i][j])}',
            va='center', ha='center', color=color)
plt.xlabel('Predicted Value')
plt.xticks(range(CLASSES))
plt.ylabel('Actual Value')
plt.yticks(range(CLASSES))
plt.colorbar()
plt.tight_layout()
plt.savefig(str(plot_name), bbox_inches='tight', pad_inches=0)
plt.close()

# Seed for consistency
np.random.seed(SEED)
# File locations
# Load best weights back up and make confusion matrices
classifier = Classifier(input_size=INPUTS)
weight_files = sorted(CLASS_MODEL_DIR.iterdir())
for weight_file in weight_files[:-1]:
    classifier.addLayer(file_name=weight_file, output=False)
    classifier.addLayer(file_name=weight_files[-1], output=True)
train_conf_title = 'Train Confusion Matrix'
makeConfMat(classifier, train_data, train_labels, CLASS_TRAIN_CONF,
    title=train_conf_title)
test_conf_title = 'Test Confusion Matrix'
makeConfMat(classifier, test_data, test_labels, CLASS_TEST_CONF,
    title=test_conf_title)
test_err = classifier.eval(test_data, test_labels)
print(f'Test error: {test_err:0.3f}')
with open(str(CLASS_TEST_LOSS), 'w') as loss_f:
    loss_f.write(f'{test_err:0.3f}')

```

A.9 test_autoencoder.py

```

#####
# Imports
#####
# Custom imports
from autoencoder import Autoencoder
from settings import *
# External imports
import numpy as np
import pathlib
import matplotlib
import matplotlib.pyplot as plt

def splitClasses(data, labels):
    '''
    '''
    split_data = list()
    for i in range(CLASSES):
        split_data.append(list())
    for d, l in zip(data, labels):
        idx = np.argmax(l)
        split_data[idx].append(d)
    return split_data

```

```

def getLossByClass(autoencoder, data, labels):
    '''
    '''
    loss = list()
    split_data = splitClasses(data, labels)
    for i, d in enumerate(split_data):
        print(f'Evaluating class {i}')
        loss.append(autoencoder.eval(d, d))
    return loss

def getSamplePoints(data, n):
    '''Get sample points from the given data set
    Parameters:
    -----
        data : np.ndarray
            Array of data points
        n : int
            Number of sample points
    Returns:
    -----
        np.ndarray
            Reduced list of data points
    '''
    indices = np.random.choice(np.arange(len(data)), 8, replace=False)
    return data[indices]

def drawSamples(autoencoder, data, num_samples, dir_name, title):
    '''Draw the output predictions and save them
    Parameters:
    -----
        autoencoder : Autoencoder
            Autoencoder for use in inference
        data : np.ndarray
            Array of data points
        num_samples : int
            Number of sample points
        dir_name : str
            Name of the directory to save the images to
        title : str
            Title of the plot
    '''
    sample_points = getSamplePoints(data, num_samples)
    for i, d in enumerate(sample_points):
        d_name = dir_name.joinpath(f'orig_{i}.png')
        matplotlib.image.imsave(str(d_name), d.reshape(28, 28, order='F'), cmap='Greys_r')
        pred = autoencoder.predict(d, one_hot=False)
        p_name = dir_name.joinpath(f'pred_{i}.png')
        matplotlib.image.imsave(str(p_name), pred.reshape(28, 28, order='F'), cmap='Greys_r')

# Seed for consistency
np.random.seed(SEED)
# Load best weights back up
autoencoder = Autoencoder(input_size=INPUTS)
weight_files = sorted(AUTO_MODEL_DIR.iterdir())
for weight_file in weight_files[:-1]:
    autoencoder.addLayer(file_name=weight_file, output=False)
autoencoder.addLayer(file_name=weight_files[-1], output=True)
# Test on all data and draw samples
test_err = autoencoder.eval(test_data, test_data)
print(f'Test loss: {test_err:0.3f}')
sample_title = 'Autoencoder Sample Outputs'
drawSamples(autoencoder, test_data, 8, AUTO_SAMPLE_DIR, sample_title)
# Graph loss by class
print('Testing train set')
train_loss = getLossByClass(autoencoder, train_data, train_labels)
print('Testing test set')
test_loss = getLossByClass(autoencoder, test_data, test_labels)
x = np.arange(len(train_loss))
plt.figure()
rect_width = 0.35
plt.bar(x-rect_width/2, train_loss, rect_width, label='Train')
plt.bar(x+rect_width/2, test_loss, rect_width, label='Test')
plt.title('Autoencoder Loss by Class')
plt.xlabel('Class')
plt.xticks(x)
plt.ylabel('Loss')
plt.grid(axis='y')
plt.gca().set_axisbelow(True)
plt.legend(loc='lower right')
plt.tight_layout()
plt.savefig(str(AUTO_BAR), bbox_inches='tight', pad_inches=0)
with open(str(AUTO_TEST_LOSS), 'w') as loss_f:

```

```
loss_f.write(f'{test_err:0.3f}')
```

A.10 features.py

```
#####
# Imports
#####
# Custom imports
from settings import *
from classifier import Classifier
from autoencoder import Autoencoder
# External imports
import numpy as np
import pathlib
import matplotlib
import matplotlib.pyplot as plt

def drawFeatures(weights, dir_name):
    '''Draw the output predictions and save them
    Parameters:
    -----
        weights : np.ndarray
            Array of weights for the neurons
        dir_name : str
            Name of the directory to save the images to
    '''
    for i, w in enumerate(weights):
        # Remove bias and normalize on [0, 1]
        w = w[1:]
        w -= w.min()
        w /= (w.max() - w.min())
        w_name = dir_name.joinpath(f'feat_{i:02d}.png')
        matplotlib.image.imsave(str(w_name), w.reshape(28, 28, order='F'),
                                cmap='Greys_r')

# Seed for consistency
np.random.seed(SEED)
# Load best weights back up for each model
autoencoder = Autoencoder(input_size=INPUTS)
for weight_file in sorted(AUTO_MODEL_DIR.iterdir()):
    autoencoder.addLayer(file_name=weight_file)
classifier = Classifier(input_size=INPUTS)
for weight_file in sorted(CLASS_MODEL_DIR.iterdir()):
    classifier.addLayer(file_name=weight_file)
# Neurons to check
neuron_count = classifier.layers[0].num_neurons
neurons = np.random.choice(np.arange(neuron_count), 20, replace=False)
class_weights = classifier.layers[0].w[neurons]
drawFeatures(class_weights, CLASS_FEAT_DIR)
auto_weights = autoencoder.layers[0].w[neurons]
drawFeatures(auto_weights, AUTO_FEAT_DIR)
```